

AI Support Agent for Helmpath

Prepared for Helmpath Inc. · February 8, 2026

Client	Helmpath Inc. (B2B logistics SaaS, ~120 customers)
Engagement	Fixed-price build + 12-month managed retainer
Timeline	7 weeks to general availability
Investment	€8,400 (build) + €620/month (managed)
Channels	In-app widget, email tier 1, Slack Connect for top accounts

Executive Summary

Helmpath's support team currently handles around 480 tickets per month with a median response time of 4.2 hours. Roughly 38% of those tickets are repeated questions about routing rules, API rate limits, and the freight-label CSV format. This proposal deploys an AI support agent that resolves those tier-1 tickets autonomously, drafts replies for tier-2, and routes anything involving billing or contract terms to a human without attempting to answer.

Understanding of Requirements

Your team has three explicit requirements that shape this build. First, the agent must cite its sources from your documentation site in every reply, since wrong information about routing rules has real shipping consequences for your customers. Second, billing and contract questions must never be answered by the AI — they must be tagged and routed to a human. Third, the agent must be able to read a customer's recent API request log when they ask 'why did my shipment fail' so the answer can be specific rather than generic.

Proposed Solution

Retrieval-augmented agent core

We index your documentation site, runbook, and the past 18 months of resolved support tickets into a vector store (pgvector on your existing Postgres instance). Every reply is grounded — if the retrieval step does not return a high-confidence chunk, the agent explicitly says so and escalates rather than fabricating.

Tool-using agent for log lookups

For 'why did this fail' questions, the agent has read-only access to a scoped subset of your API request log: the customer's own tenant, last 7 days, request and response status only. It cannot see request bodies or PII. This lets it answer specifically without introducing data-handling risk.

Tier-1 / tier-2 / human routing

An intent classifier sits in front of the agent. Tier-1 (clear documentation question) gets answered directly. Tier-2 (ambiguous or partially-covered) gets a draft reply written and queued for a human reviewer in your existing Front instance. Tier-3 (billing, contracts, security questions, anything mentioning a competitor) is routed to a human with no draft.

Feature Breakdown

Feature	Priority	Notes
In-app widget on app.helmpath.com	Must	Custom Tailwind component, matches your design system
Email triage and auto-reply	Must	Front integration
Cited replies (every claim links to a doc)	Must	Hard requirement — no uncited generations
API log lookup tool	Must	Scoped read-only, audited
Slack Connect support for enterprise accounts	Should	Top 12 accounts, opt-in
Conversation analytics dashboard	Should	Resolution rate, escalation rate, CSAT proxy
Multilingual (DE, FR)	Nice	Phase 2

Project Timeline

Week	Milestone	Deliverable
1	Discovery & corpus build	Indexed documentation, 18-month ticket history
2	Agent core	Working retrieval-augmented replies in staging
3	Tool layer	API log read access, audit logging
4	Routing & guardrails	Intent classifier, billing/contract escalation tested
5	In-app widget	Live on staging app.helmpath.com
6	Internal pilot	Helmpath staff use the agent on real tickets, daily review
7	GA rollout	Live to all customers, dashboard, runbook delivered

Pricing

Component	Cost
Discovery, corpus build, agent core	€3,800
Tool layer, routing, guardrails	€2,400
In-app widget, Front integration, dashboard	€1,800
Pilot week + GA rollout + training	€400
Managed retainer (12-month minimum)	€620 / month

Relevant Experience

We have shipped two comparable B2B SaaS support agents in the past year: an HR-tech vendor (~85 customers, 31% deflection rate achieved by month two) and a developer-tooling startup (~210 customers, 44% deflection, used as a reference customer). The logistics domain is new to us but the architecture transfers directly, and we will spend the first week of discovery with your support team learning the failure modes that matter most to your customers.

Why Nexrai Labs

- Citation-grounded agent — every claim is linked to a source in your docs.
- Strict no-go zones for billing and contract questions, enforced by classifier and not by prompt-only.
- Senior-only delivery team, no offshore subcontracting.
- Sits on top of your existing Front and Postgres infrastructure.
- Transparent retainer with logged hours and quarterly reviews.

Appendix

Architecture diagram, sample cited replies, and a draft of the security review document (covering log scoping, audit trail, and prompt-injection mitigations) are included in the shared folder.